

Reliability-Screened Routing Under Predictable and Stochastic Travel Variation: Matched Comparisons of Conservative Planning and Statistical Screening

Suhan Ali

Independent researcher, Lahore, Pakistan

Corresponding author: syedsuhanali9@gmail.com ORCID: 0009-0004-5079-7684

Abstract

Routes planned on nominal travel times can fail under predictable congestion and random delay. This paper tests a two-layer procedure for the capacitated vehicle-routing problem with time windows. A replaceable solver proposes plans, and internal window contraction adds time buffers without changing customer promises. A plan is admitted only if a one-sided 95% Wilson lower bound on the probability that at least 95% of customers are served on time reaches 0.95, and the selected plan is checked once on fresh scenarios. On 53 synthetic instances of 20–200 customers with OR-Tools candidates, the screened procedure exceeded uniform travel-time padding by 8.3–21.4 percentage points. Capping each customer’s contraction at its direct-reachability limit raised admission from 32 to 42 instances, at 2.5% more distance. Two confirmatory studies followed, each on 53 new instances from the same instance generator under a publicly timestamped protocol. The first confirmed all four prespecified hypotheses: screening exceeded conservative padding by 15.7 points, and capping added 11.7. The second replaced OR-Tools with hybrid genetic search (HGS): screening with capping exceeded HGS conservative padding by 27.9 points and differed from the OR-Tools procedure by -0.2 points (95% interval -4.0 to 3.7). A conservative-speed fallback suggested post hoc helped the uncapped procedure but did not measurably help the capped one. These results use the original objective weights; with equal weights, the advantage over padding is not distinguishable from zero. Pooled service stays below the 95% target, mainly through fallback plans. Results depend on assumed disturbance parameters and synthetic geometry.

Keywords: vehicle routing; time windows; stochastic travel times; chance constraints; simulation-based validation; reproducibility

MSC Classification: 90B06, 90C59, 90C15, 62G09

1 Introduction

The vehicle-routing problem with time windows (VRPTW) asks a fleet to serve customers within agreed time windows, subject to vehicle capacity and shift length [19, 21]. The standard formulation plans with one fixed matrix of travel times. In practice, travel times change with the time of day and with random delays, and a delay early in a route can push later customers past their windows. A plan that minimizes nominal distance takes no account of this risk.

One simple safeguard is to plan with a slower speed than expected, which pads every trip. This paper examines another approach. A standard VRPTW solver proposes plans with extra internal slack, each plan is tested on simulated days, and a plan is admitted only if it meets a statistical reliability target; otherwise a prespecified fallback plan is used. The experiments ask whether this screened procedure gives more reliable service than uniform padding on the same instances, whether any gain comes from route structure or simply from using more vehicles, and what it costs in distance, driver time and fleet size. They also examine how the choice of candidate generator, knowledge of the daily congestion profile and a limit on each customer’s buffer change how often a plan passes the test and how reliable it is out of sample. Finally, they test whether the main findings hold on new instances under prespecified protocols, whether they survive a change to a state-of-the-art solver, and whether a conservative-speed plan is a better fallback when no plan passes. An exploratory analysis examines how sensitive the findings are to the objective weights.

The procedure keeps *candidate generation* separate from *stochastic validation*. The solver is a nominal VRPTW solver; it is not presented as an exact method for a time-dependent stochastic program. Candidate routes are tested after they are solved, first on a screening bank of simulated scenarios and then once on an untouched validation bank. The main reliability event allows a small share of customers to be late on a given day. The probability that every customer is on time falls steadily as the number of customers grows, even for good plans, so it is a poor target for large instances.

The individual components are established: time-dependent speed profiles, multi-criteria routing, chance constraints, Monte Carlo evaluation and Wilson intervals. The contribution is the way they are combined and the evidence for it. The routing solver is treated as a replaceable candidate generator, and every reliability decision is made after it. Internal window contraction produces a family of candidate plans, a statistical screen admits or rejects each one, and an independent scenario bank validates the plan chosen. Capping each customer’s contraction at its direct-reachability limit removes one cause of failure, in which contraction alone makes a customer unreachable even by a direct trip. The candidate family and the selection rule work together, so a comparison with a single uncontracted plan cannot attribute the gain to the selection rule alone.

The evidence comes in three parts. A development study on 53 synthetic instances shaped the design. It compares the procedure with uniform travel-time padding and with ablations that isolate clock-profile information, fleet size, contracted candidates and screening, and its findings are exploratory. A first confirmatory study then tested four prespecified hypotheses on 53 new instances, and an exploratory analysis varied the objective weights. A second confirmatory study, on another 53 new instances, replaced OR-Tools with hybrid genetic search [23, 25] and tested the conservative-speed fallback that the first study had suggested post hoc. Both confirmatory protocols were publicly timestamped before their instances were solved.

2 Related work

Solomon’s benchmark study established the standard VRPTW experimental setting [19]; Toth and Vigo provide a broad treatment of vehicle-routing models and algorithms [21]. Time-dependent routing requires care because a naive multiplier applied at departure can violate

first-in-first-out (FIFO). Ichoua, Gendreau, and Potvin introduced the integrated speed-profile construction used here [12]; Ghiani and Guerriero later analysed its representational properties [11].

Stochastic VRPTW research has treated uncertain travel and service times using chance constraints, recourse, approximate route-feasibility tests, and data-driven formulations [6, 9, 14, 18]. Ehmke, Campbell, and Urban embed per-customer service levels in a routing method [8]; Taş et al. jointly model time dependence, stochastic travel, service costs, and soft windows [20]; and Bomboi, Buchheim, and Prunte develop route-feasibility checks for correlated and time-dependent travel times [6]. Vareias, Repoussis, and Tarantilis optimize self-imposed time windows under stochastic travel times [22], which is the closest conceptual analogue to the internal window contraction used here.

A second line of work builds the uncertainty model into the optimization. Agra et al. solve a robust VRPTW in which travel times range over an uncertainty set [1]. Jaillet, Qi and Sim route to minimize a requirements violation index, which measures how far arrival times can breach their windows when the travel-time distribution is only partly known [13]. Zhang et al. develop a riskiness index for time windows under uncertainty, together with an exact solution algorithm [26]. Gendreau, Jabali and Rei survey stochastic routing models and the directions that remain open [10]. These methods require a solver built for their uncertainty model, and none of them is compared with the procedure in this paper.

The procedure in this paper differs from both lines of work in where the statistical claim is made. It does not build a stochastic feasibility approximation into a specialist route-search algorithm. It takes routes from a replaceable nominal solver, searches a family of internal buffers on one scenario bank, and checks the chosen plan once on a held-out bank. Its service event limits the share of customers served late across the whole day, and no per-customer chance constraints are used. This solver-agnostic design is narrower than a new stochastic optimizer, but it makes selection failures and out-of-sample reliability directly visible. Monte Carlo approximations of chance constraints need a separate check of candidate feasibility [2]. The Wilson bound is used because normal-approximation intervals perform poorly for probabilities close to one [7, 24].

The main candidate generator is Google OR-Tools [16], a practical local-search framework with no optimality guarantee. PyVRP implements a much stronger hybrid genetic search (HGS) for standard VRPTW variants [23, 25]. Here it serves as a deterministic benchmark and, in the second confirmatory study, as an alternative candidate generator.

3 Problem definition

Let $N = \{0, 1, \dots, n\}$ be the nodes, with depot 0 and customers $C = N \setminus \{0\}$. Let V be a homogeneous fleet. Customer i has demand δ_i , deterministic service time s_i , and offered service-start window $[l_i, u_i]$. Vehicle capacity is Q , and the maximum shift duration is $H^{\max}$. The nominal distance and planning time on arc (i, j) are d_{ij} and $\bar{t}_{ij}$.

The planning variables are $x_{ijv} \in \{0, 1\}$ for arc use, $z_{iv} \in \{0, 1\}$ for customer assignment, $y_v \in \{0, 1\}$ for vehicle use, service-start time a_{iv} , depot departure d_v , depot return r_v , and slack deficit $G_{iv} \geq 0$. The assignment variables give each customer one vehicle and activate the corresponding customer-vehicle copy:

$$\sum_{v \in V} z_{iv} = 1 \quad i \in C, \quad (1)$$

$$\sum_{j \in N \setminus \{i\}} x_{ijv} = \sum_{j \in N \setminus \{i\}} x_{jiv} = z_{iv} \quad i \in C, v \in V, \quad (2)$$

$$\sum_{j \in C} x_{0jv} = \sum_{i \in C} x_{i0v} = y_v \quad v \in V. \quad (3)$$

Because demand is delivery-only, route capacity needs no cumulative-load variables:

$$\sum_{i \in C} \delta_i z_{iv} \leq Q y_v \quad v \in V. \quad (4)$$

The time-propagation constraints below rule out customer-only subtours: summing them around any directed customer cycle would give $0 \geq \sum_i (s_i + \bar{t}_{ij}) > 0$. The argument does not need positive demands. Time windows apply only to assigned copies:

$$l_i z_{iv} \leq a_{iv} \leq u_i z_{iv} \quad i \in C, v \in V. \quad (5)$$

Separate departure and return variables avoid using one time variable for two depot visits. With a valid global deactivation constant $M \geq H^{\max} + \max_i s_i + \max_{ij} \bar{t}_{ij}$,

$$a_{jv} \geq d_v + \bar{t}_{0j} - M(1 - x_{0jv}), \quad (6)$$

$$a_{jv} \geq a_{iv} + s_i + \bar{t}_{ij} - M(1 - x_{ijv}), \quad (7)$$

$$r_v \geq a_{iv} + s_i + \bar{t}_{i0} - M(1 - x_{i0v}). \quad (8)$$

The bounds $0 \leq d_v, r_v \leq H^{\max} y_v$ deactivate unused vehicles, and $r_v - d_v \leq H^{\max} y_v$. The code enforces these timing relations through OR-Tools routing dimensions and never builds (8) explicitly, so no big- M relaxation or big- M diagnostic is claimed for the code.

For target slack s^* ,

$$G_{iv} \geq a_{iv} - (u_i - s^*) z_{iv}, \quad (9)$$

$$G_{iv} \geq 0. \quad (10)$$

At the optimum, $G_{iv} = \max\{0, s^* - (u_i - a_{iv})\}$ when $z_{iv} = 1$, and $G_{iv} = 0$ when the customer is not assigned to vehicle v .

4 Candidate objectives and scaling

The four raw criteria are

$$F_1 = c_f \sum_{v \in V} \sum_{i,j \in N} d_{ij} x_{ijv}, \quad F_2 = \sum_{v \in V} \sum_{i \in C} G_{iv}, \quad (11)$$

$$F_3 = \sum_{v \in V} (r_v - d_v), \quad F_4 = \sum_{v \in V} y_v. \quad (12)$$

F_3 counts driving, service and waiting time, and each vehicle may leave the depot when it suits the schedule, so no vehicle is charged for waiting at the depot before its first service can start.

Reference values come from four heuristic solves, each minimizing one criterion. For each criterion k , L_k and U_k are the smallest and largest values among these anchor solutions. They are *scaling ranges*, not utopia and nadir points, and need not bound the Pareto frontier. No value is clipped. The normalized score is

$$Z(\mathbf{x}) = \sum_{k=1}^4 w_k \frac{F_k(\mathbf{x}) - L_k}{\max\{U_k - L_k, 1\}}. \quad (13)$$

The floor of 1 in the denominator avoids division by zero when a discrete criterion takes the same value in every anchor solution. The baseline weights $\mathbf{w} = (0.5693, 0.2643, 0.1055, 0.0609)$ are the principal-eigenvector weights of one Analytic Hierarchy Process (AHP) pairwise-comparison matrix with consistency ratio 0.0254 [17]. Appendix A and `code/ahp_weights.py` give the matrix and the calculation. The matrix was set for an earlier list of criteria in which the third

criterion was total distance, and its third row and column were not re-elicited when that criterion became driver duration. The weights are therefore a fixed experimental preference vector that defines the candidate-selection objective. They were fixed before any experiment reported here, were not elicited from a fleet operator, and make no claim about the preferences of any industry. Section 10.8 reports a sensitivity analysis over four weight vectors.

In OR-Tools, distance is an arc cost, driver duration is a span cost on each vehicle’s time dimension, slack deficit is a soft upper bound at each customer, and fleet use is a fixed cost per vehicle. Integer coefficients are scaled once per criterion. The main planning matrix is nominal and fixed, so the original candidate generator does not use clock-dependent travel times. The *conservative-speed* baseline replaces 40 km/h with 33.3 km/h throughout the planning matrix. This pads every trip by 20%, including off-peak trips, so the comparison does not isolate the value of knowing when the peak occurs.

4.1 Clock-aware approximation and deterministic guarantee

A separate distance heuristic uses the time of day. For each origin, it freezes the departure time estimated from the previous schedule, integrates (14) at that time to build the origin’s outgoing travel-time row, and solves again. Two passes receive 2.5 seconds each. The first reference schedule is the retained nominal plan, whose solve cost has already been incurred separately. For a customer, the departure estimate is the service start plus the service time. For a depot arc, it is the previous departure for the first customer of each route and the mean of the used depot departures for other customers. Each resulting route sequence is rescheduled with the exact deterministic FIFO profile, and binary search finds its latest feasible whole-minute depot departure. Routes that fail even when leaving at time zero are rejected. The shortest deterministically feasible output is kept, with duration breaking ties. The method is an approximation that reports its failures; it is not a departure-dependent callback or an exact optimizer for the time-dependent VRPTW.

For a feasible conservative-speed plan, the evaluation speed at $\sigma = 0$ never drops below the planning speed during the planned shift. Starting from the same depot departure, FIFO and induction along each route show that the evaluated service starts and the return happen no later than planned. At $\sigma = 0$, such a plan serves every customer on time by construction. The deterministic feasibility of the clock-aware approximation is also checked before a plan is accepted. Neither property is an experimental reliability gain.

5 Time-dependent stochastic evaluation

5.1 FIFO travel times

Partition the day into periods $[T_k, T_{k+1})$ with strictly positive speeds ν_k . The deterministic time-dependent travel time $t_{ij}(\tau)$ is the smallest $\Delta \geq 0$ satisfying

$$\int_{\tau}^{\tau+\Delta} \nu(u) du = d_{ij}, \quad \nu(u) = \nu_k \text{ for } u \in [T_k, T_{k+1}). \quad (14)$$

Because the distance covered grows strictly with time, the arrival map $\tau \mapsto \tau + t_{ij}(\tau)$ is nondecreasing, so FIFO holds.

5.2 Correlated disturbances

For scenario ω , arc travel is

$$T_{ij}(\tau, \omega) = t_{ij}(\tau) \Omega_{ij}(\omega), \quad (15)$$

where

$$\Omega_{ij}(\omega) = \exp\{\varepsilon_{ij}(\omega)\}, \quad (16)$$

$$\varepsilon_{ij}(\omega) = \sigma \left(\sqrt{\rho} \xi_0(\omega) + \sqrt{1 - \rho} \xi_{ij}(\omega) \right). \quad (17)$$

Here ρ is the correlation between the latent Gaussian log-errors. It is not the Pearson correlation between the multipliers, which is

$$\text{Corr}(\Omega_i, \Omega_j) = \frac{\exp(\sigma^2 \rho) - 1}{\exp(\sigma^2) - 1}. \quad (18)$$

The multiplier has median 1 and mean $\exp(\sigma^2/2)$, so the deterministic profile gives the median travel time, which is slightly below the mean.

Scenario banks do not depend on the order in which routes are evaluated. A common factor is drawn once per scenario, and each directed arc gets its own pseudorandom stream keyed by the bank seed and the arc's endpoints. Plans evaluated on the same bank see identical disturbances on the arcs they share.

6 Screening and independent validation

For a fixed plan R and scenario ω , let

$$q(R, \omega) = \frac{1}{|C|} \sum_{i \in C} \mathbf{1}\{a_i(R, \omega) \leq u_i\} \quad (19)$$

be the on-time customer proportion. The primary service event is

$$J_\kappa(R, \omega) = \mathbf{1}\{q(R, \omega) \geq \kappa\}, \quad \kappa = 0.95, \quad (20)$$

and the target is

$$\Pr[J_\kappa(R, \omega) = 1] \geq \alpha, \quad \alpha = 0.95. \quad (21)$$

The stricter probability that no customer in the fleet is late on a given day is also reported. It is not used for admission, because it falls mechanically as the number of customers grows.

For N scenarios and empirical success $\hat{p}$, the one-sided Wilson lower bound at confidence $1 - \eta$ is

$$L(\hat{p}) = \frac{\hat{p} + z^2/(2N) - z\sqrt{\hat{p}(1 - \hat{p})/N + z^2/(4N^2)}}{1 + z^2/N}, \quad z = z_{1-\eta}. \quad (22)$$

Candidate generation contracts only the internal planning deadline, to $[l_i, \max\{l_i, u_i - \beta\}]$. The offered window $[l_i, u_i]$ does not change, and every scenario is scored against the original u_i . For each buffer β , a candidate is generated with the slack-aware objective and assessed on the screening bank. The anchor ranges come from the uncontracted ($\beta = 0$) instance and stay fixed across the buffer search. Among the candidates whose screening lower bound reaches α , the one with the smallest normalized planned score (13) is selected. In simulation every vehicle leaves at its planned d_v ; departure times do not change with the scenario. The selected plan is then evaluated once on a larger, untouched validation bank, so the search never selects on the evidence used for the final report.

A prespecified fallback makes the procedure's output defined on every instance. If no candidate passes screening, the procedure returns the uncontracted slack-aware plan ($\beta = 0$). Validation never changes this choice; it only estimates the performance of the plan already chosen. Results for selected plans alone are always labelled as such, and the main procedure comparison includes all instances.

Three reliability quantities appear in the results, and each answers a different question:

- **Admission criterion.** The Wilson lower bound (22) on the screening bank, computed for each candidate. A candidate is admitted when it reaches $\alpha = 0.95$.
- **Selected-plan probability.** The validated estimate of $\Pr[J_\kappa = 1]$ for a plan that passed screening, from the untouched bank.
- **Procedure-level probability.** The validated probability averaged over all instances, including those on which screening admitted nothing and the fallback was returned.

The chance constraint (21) is what screening tests for each admitted plan. The procedure-level average falls below α whenever fallback plans are weaker, so the results report admitted plans and the whole procedure separately.

6.1 Generator replacement and reachability-capped contraction

Screened conservative speed uses the same buffer grid, admission bank and Wilson threshold, but its candidates minimize distance on the padded matrix. Among the passing plans it chooses the shortest, breaking ties by planned duration and then by buffer. Its fallback is the uncontracted conservative plan, or the retained uncontracted slack-aware plan if no conservative plan was found. The procedure’s output is therefore defined on all 53 instances, while the conservative generator itself still finds a plan on 46 of them.

Let $e_i = \max\{l_i, \bar{t}_{0i}\}$ be the earliest direct service time under the nominal matrix. The capped weighted variant uses the internal upper bound

$$u_i^{\text{cap}}(\beta) = \min\{u_i, \max\{l_i, u_i - \beta, e_i\}\}. \quad (23)$$

This stops contraction alone from making a customer unreachable by a direct trip when it was reachable before. It does not guarantee that a joint routing is feasible or that the admission target is met. When a contracted model is identical to one already solved, its archived candidate is reused; only changed models are solved again. Selection still uses the fixed original scaling ranges, the weighted score and the same fallback.

To diagnose failures, a shortest-path relaxation uses arc cost $\bar{t}_{ij} + s_i$ (with $s_0 = 0$) and allows intermediate customers without capacity or window restrictions. Its depot-to- i distance is a valid lower bound even if rounded distances violate the triangle inequality. A contracted deadline below this bound certifies infeasibility. When there is no certificate, a missing solver output is labelled *no feasible plan found*: reaching the time limit without a plan does not prove infeasibility, and passing the reachability test does not prove feasibility. Solver status information was not kept in the archived runs and is not inferred afterwards.

7 Experimental design

Synthetic instances place a central depot and uniformly distributed customers on a 100×100 grid. Demand is uniform on 10–49 units, service time is 10 minutes, window width is uniform on 90–149 minutes, capacity is 1000, and the shift limit is 600 minutes. Time zero is 06:00 local time, so clock time is the relative schedule time plus 360 minutes. Table 1 gives the seeded instance sizes, and Table 2 lists every model and simulation constant.

Eight configurations are compared:

1. **Nominal distance:** fixed 40 km/h planning matrix and a distance-minimizing objective.
2. **Conservative-speed distance:** distance minimization with a uniform 33.3 km/h planning matrix.
3. **Slack-aware:** objective (13) with $\beta = 0$, which shows the effect of the objective without screening.

Table 1 Seeded synthetic experimental design.

Customers	Vehicles available	Instances	Seeds
20	6	20	7200–7219
50	14	15	7500–7514
100	26	10	8000–8009
200	40	8	9000–9007

Table 2 Planning, disturbance, and statistical parameters.

Parameter	Retained value
Nominal/free-flow speed $\bar{v}$	40 km/h
Shift and clock origin	600 min, beginning 06:00 local time
Capacity; service time; target slack s^*	1000; 10 min; 15 min
Distance cost c_f	15 cost units/km
AHP weights (w_1, w_2, w_3, w_4)	(0.5693, 0.2643, 0.1055, 0.0609)
Integer objective scale	100,000
Planned-shift speed divisors	0.9 (06:00–07:00), 1.2 (07:00–09:00), and 1.0 (09:00–16:00); speed is $\bar{v}$ divided by the divisor
Outside-shift continuation	1.0 to 17:00, 1.3 (17:00–19:00), 0.9 thereafter; used if a stochastic route overruns the shift
Log-error scale σ ; latent correlation ρ	0.20; 0.50
Exploratory disturbance grid	$\sigma \in \{0.10, 0.20, 0.30\}$; $\rho \in \{0, 0.5, 0.9\}$
Service fraction κ ; target α	0.95; 0.95
Wilson tail probability η	0.05 (one-sided 95% lower bound)
Screening; validation; sensitivity scenarios	1000; 5000; 1000 per cell
Buffer grid	$\beta \in \{0, 5, \dots, 60\}$ min

4. **Fleet-matched distance:** distance minimization forced to use exactly as many vehicles as the slack-aware plan, starting from its feasible routes.
5. **Original procedure:** the complete weighted internal-buffer search, with selection on 1000 screening scenarios, the prespecified $\beta = 0$ fallback, and assessment on 5000 fresh scenarios.
6. **Screened conservative:** conservative-speed distance candidates with the same admission test and the fallback defined above.
7. **Capped procedure:** the original weighted procedure with customer-specific reachability caps.
8. **Clock approximation:** two matrix updates from estimated departure times, followed by exact deterministic rescheduling.

The nominal shift ends at 16:00. The full-day speed profile continues past that time, because lognormal delays are unbounded and a stochastic route can overrun. Evening periods lie outside the planned shift but stay in the evaluator. Fleet-matched distance starts from the slack-aware routes. Because this could carry over their structure, the starting point is treated as part of that comparator’s definition.

The screening grid is $\beta \in \{0, 5, \dots, 60\}$ minutes. The four single-criterion anchor solves get two seconds each to stabilize the scaling ranges, and each candidate solve gets a five-second OR-Tools budget. If a candidate solve returns no feasible plan, the same five-second call is repeated at most twice more. The first feasible result is kept; there is no best-of-several selection.

A new model with a reachability certificate is not solved at all. The clock approximation splits one five-second budget between its two solves. Six experiment workers ran in parallel, and elapsed times and software versions are recorded. Within an instance, every configuration uses the same fresh validation bank.

The reliability outcomes are service-event probability, zero-lateness fleet-day probability, mean on-time fraction, expected number of late customers, mean and 95th-percentile total lateness, and mean route success. The operational outcomes are distance, driver duration and fleet use. Simulated duration under the common evaluation model is also reported, to separate actual route time from each generator’s planned F_3 . Differences are paired by instance, and intervals across instances are kept separate from binomial Monte Carlo intervals within an instance. Across-instance intervals use a paired nonparametric bootstrap with 20,000 resamples and a fixed analysis seed. A claim that distances are equivalent would need a prespecified equivalence margin, so failing to reject a difference is not treated as proof of equality.

8 First confirmatory study and weight-sensitivity design

A confirmatory protocol was fixed before any new instance was generated or solved. It sets out four directional hypotheses, their tests and decision rules, a failure policy, and 53 instance seeds equal to the development seeds shifted by 60000. The protocol, the code that runs and analyses it, and a SHA-256 manifest of those files were committed to a public repository (commit 41fae67d9520bd5dd01fcac87e680a5a69325d10) and archived on Zenodo [3] before the study was run. The analysis program will not run unless the complete prescribed study is present and every frozen file matches the manifest.

Each hypothesis concerns validated service-event probability on the fresh bank, with the instance as the unit. H1 (primary): on instances where conservative-speed planning returns a plan, the original procedure exceeds it. H2 (primary): on all instances, the capped procedure exceeds the original procedure. H3: screening admits the capped procedure on more instances than the original procedure, tested with an exact McNemar test at the 0.05 level. H4: on the instances of H1, the advantage of the original procedure over conservative speed is larger at $\sigma = 0.30$ than at $\sigma = 0.10$ with $\rho = 0$. H1 and H2 use 97.5% intervals, a Bonferroni adjustment for two primary hypotheses, and each is confirmed only if its paired bootstrap interval lies above zero. H4 uses a 95% interval. H3 and H4 are secondary: they carry less weight than H1 and H2 and cannot rescue an unconfirmed primary hypothesis. Holm’s procedure in place of the Bonferroni adjustment would not change either primary decision. H4 compares fixed plans: each plan is screened under the primary setting ($\sigma = 0.20$, $\rho = 0.5$) and then evaluated across the sweep without re-screening.

A configuration without a plan scores zero on that instance, in validation and in every sweep cell, so a failure counts against the configuration that failed. A software error is recorded with its traceback, is never scored as routing performance, and blocks the analysis.

Parameters, budgets and scenario counts are those of the development study (Table 2). The study ran on an Apple M1 Pro laptop with ten cores under macOS 26.3.1, with Python 3.12.15, NumPy 2.5.3 and OR-Tools 9.15.6755; the development runs used Python 3.12.14 on Linux. Stage 1 ran the four instance sizes as four parallel processes, where the written protocol lists a sequential loop, and stage 2 used six workers as specified. These are the deviations from the written procedure. The scheduling change affects only the CPU time available to each time-limited solve.

The weight analysis is exploratory. The original procedure was rerun on the development instances with four weight vectors: the original vector, the original with its two largest weights swapped, a distance-heavy vector and equal weights (Section 10.8). Each rerun regenerates the candidate family, screens it on the same bank and validates the selected plan or fallback on the same fresh bank, with the same five-second budget. The vectors were chosen before the runs,

and each was run once.

The benchmark check is also exploratory. It uses the eight R1 and RC1 Solomon files available in the ML4VRP 2024 competition repository [15], which hosts Solomon’s instances [19]. The C1 instances are left out because their horizon of 1236 time units is longer than one working shift. Coordinates are read as kilometres, with the study’s rounded Euclidean distances and a planning speed of 40 km/h, so one Solomon time unit equals 1.5 minutes. Ready and due times are multiplied by 1.5 and rounded down, service time becomes 15 minutes, and the capacity of 200 and the 25 vehicles are kept. The shift length is the scaled depot due time, 345 minutes for R1 and 360 minutes for RC1, and the shift starts at 06:00 so that the morning peak falls inside it. Planning times are rounded up to whole minutes, as in the main study. This mapping keeps Solomon’s relation between travel time and window width, because one distance unit takes one Solomon time unit to drive at the planning speed. The C1 instances could be fitted by compressing time, but that would change how tight their windows are relative to travel times, so the eight instances are a convenience sample of the R1 and RC1 types. The pipeline code is unchanged; only the instance loader and the parameter constructor are replaced.

9 Second confirmatory study design

The first confirmatory study left two questions open. Its candidate generator produces weaker routes than the state of the art, so the benefit of screening might shrink with candidates from a stronger search. Its post hoc analysis also suggested that returning the conservative-speed plan when screening admits nothing would raise pooled service (Section 10.7), but that rule was chosen after the results were seen. A second protocol tested both questions on 53 further instances, with seeds equal to the development seeds shifted by 80000.

The second generator is the hybrid genetic search of PyVRP [23, 25], version 0.14.0. It solves the same planning model as the OR-Tools generator: the static matrix at the planning speed, hard and possibly contracted windows, capacity, shift length and fleet size, with each vehicle leaving at the latest feasible time. Its weighted objective keeps the distance, driver-duration and fleet terms with the same coefficients. PyVRP cannot express the slack-deficit term, so that term is left out of generation and kept in candidate selection. Each solve stops after 10,000 iterations and uses the instance seed as its random seed, so each HGS solve is deterministic for given inputs. Both generators share the scaling ranges from the unchanged OR-Tools anchor solves. Those solves are time-limited, so the HGS results are deterministic only once the ranges are fixed. The contraction grid, capped windows, reachability certificates, scenario banks, admission rule and selection score are the same as in the OR-Tools procedure. HGS conservative-speed planning minimizes distance at 33.3 km/h. The OR-Tools half reruns the frozen pipeline of the first study unchanged.

Two fallback rules are compared. The slack-aware rule is the one used in the first study; for HGS it returns the uncontracted weighted plan. The conservative rule returns the same generator’s conservative-speed plan when screening admits nothing, or the slack-aware plan when there is no conservative plan. Neither rule changes a screening decision or uses a validation outcome.

The protocol fixes four hypotheses on validated service-event probability, with the instance as the unit. G1 (primary): on instances where HGS conservative-speed planning returns a plan, the HGS capped procedure with conservative fallback exceeds it. G2 (primary): on all instances, the conservative fallback improves the OR-Tools capped procedure. G3 (secondary): the same holds for the uncapped OR-Tools procedure. G4: the difference between the HGS and OR-Tools capped procedures, both with conservative fallback, is estimated with a 95% interval and no directional decision. G1 and G2 use 97.5% intervals and G3 a 95% interval, with the bootstrap, failure policy and decision rules of the first study. The protocol requires the abstract and conclusion to report any unconfirmed primary hypothesis.

The protocol, the HGS code, the analysis program and a SHA-256 manifest covering them and the frozen files of the first study were committed to the public repository (commit 3ba6b71288a2a51872ab6a877fa86eb28ff7671d) and archived on Zenodo [4] before any instance of the study was solved. Before the runs began, a gate script checked the manifest, the pushed commit and tag, and that no results existed yet. An earlier draft of the protocol used seeds shifted by 70000. The OR-Tools half was run on those seeds before that draft was timestamped, and its console output, which lists the instances that screening admitted, was seen. That seed set was abandoned in full. Its records are kept in the archive but enter no analysis, the protocol discloses them, and the study used the seeds shifted by 80000 with no other change. The HGS generator was built and tested on the development instances, so its development results are exploratory. The study ran on the same machine and software as the first confirmatory study, with PyVRP 0.14.0, and followed the written procedure without deviation.

10 Results

10.1 Complete procedures on matched instances

Table 3 compares every original configuration on the same subset of instances, those on which conservative-speed planning returned a plan. The paired contrasts are computed from differences within each instance, so their intervals measure the comparison directly; they are not obtained by subtracting the endpoints of separate intervals.

Table 3 Validated service-event probability (%) on the same 46 instances: every entry in a row has the displayed denominator. Entries are mean [95% instance-bootstrap interval]. N, C, S, F and P denote nominal, conservative-speed, slack-aware, fleet-matched and original procedure, respectively.

n	k	N	C	S	F	P
20	18	21.5 [18.8, 24.4]	64.2 [61.1, 67.4]	53.3 [47.6, 59.4]	21.7 [18.9, 24.8]	85.7 [74.4, 95.0]
50	13	17.7 [16.1, 19.6]	59.0 [54.6, 63.1]	45.8 [43.1, 48.4]	16.9 [14.8, 18.9]	76.8 [62.3, 89.7]
100	8	25.4 [23.5, 27.1]	71.6 [69.1, 74.0]	55.3 [51.5, 59.3]	26.2 [23.6, 29.0]	87.6 [74.2, 97.5]
200	7	26.7 [24.9, 28.8]	74.6 [72.2, 77.0]	61.7 [58.3, 66.8]	27.7 [25.5, 30.6]	82.9 [69.3, 94.3]

Table 4 Paired differences in service probability (percentage points) on the same subsets as Table 3. SC denotes screened conservative speed with its prespecified fallback.

n	S minus C	P minus C	SC minus C
20	-11.0 [-16.1, -6.0]	21.4 [9.6, 31.3]	28.2 [21.4, 34.3]
50	-13.1 [-17.4, -8.9]	17.9 [2.8, 32.0]	11.8 [2.7, 22.2]
100	-16.2 [-20.6, -11.7]	16.0 [4.8, 24.8]	11.2 [2.8, 19.4]
200	-12.9 [-18.5, -5.7]	8.3 [-5.1, 20.0]	5.7 [0.0, 12.9]

In the paired contrasts of Table 4, the complete original procedure exceeds conservative-speed planning by 8.3–21.4 percentage points, depending on size. The slack-aware objective without screening shows no such advantage. At $n = 200$ the interval includes zero, so the positive mean at that size is not conclusive.

10.2 Resources on the common subset

Scheduled driver duration differs from the realized driving, service and waiting time in simulation: the conservative planner budgets padded arc times, while in simulation those arcs may be driven faster. Table 5 therefore reports both. Absolute distance premiums also depend on solver quality:

Table 5 Paired original-procedure changes relative to conservative speed on the common subset. Distance and duration are mean percentage changes; fleet is an absolute difference. Scheduled duration is each planner’s own F_3 . Simulated duration uses the same realized traffic model and planned departures for both configurations.

n	Distance	Scheduled time	Simulated time	Fleet
20	4.9 [2.4, 7.3]	0.3 [−3.8, 4.8]	8.5 [4.4, 12.9]	0.22 [0.00, 0.44]
50	2.3 [−0.6, 5.2]	−7.5 [−12.2, −2.7]	−2.0 [−7.0, 2.8]	−0.46 [−1.00, 0.00]
100	2.0 [−1.3, 5.1]	−3.0 [−9.4, 2.5]	2.2 [−4.3, 7.6]	0.12 [−0.38, 0.62]
200	2.4 [−1.4, 5.7]	−7.7 [−15.2, −1.0]	−2.8 [−10.6, 4.1]	−0.71 [−2.57, 1.00]

in the five-second deterministic benchmark, OR-Tools routes are about 11% longer than HGS routes at $n = 200$.

The original procedure does not use fewer resources at every size. At $n = 20$ it has slightly more scheduled time and more vehicles, and at $n = 100$ it also uses slightly more vehicles. The simulated-time interval excludes zero only for the increase at $n = 20$. The service gain is therefore bought with some extra operating resources.

10.3 Screening a replaceable generator and capping contractions

Table 6 Procedure outcomes on all instances. Selected counts instances where the lower bound of some candidate on the screening bank reached 0.95; retained counts the selections whose lower bound on the fresh validation bank also reached 0.95. Probability includes fallback plans as well as selected plans.

n	Procedure	Selected / instances	Retained / selected	Probability (%)
20	Original procedure	14/20	14/14	81.6 [70.2, 91.9]
20	Screened conservative	15/20	15/15	87.7 [79.0, 94.7]
20	Capped procedure	14/20	14/14	81.6 [70.3, 92.0]
50	Original procedure	8/15	8/8	72.7 [59.2, 86.4]
50	Screened conservative	4/15	4/4	67.5 [58.2, 77.3]
50	Capped procedure	10/15	10/10	80.1 [66.9, 90.8]
100	Original procedure	6/10	5/6	81.3 [68.8, 93.3]
100	Screened conservative	4/10	4/4	77.5 [67.2, 88.2]
100	Capped procedure	10/10	9/10	97.5 [96.7, 98.2]
200	Original procedure	4/8	3/4	80.6 [68.7, 92.3]
200	Screened conservative	2/8	2/2	78.3 [71.0, 86.5]
200	Capped procedure	8/8	7/8	97.6 [96.9, 98.2]

In Table 6, screened conservative speed selects a plan on 25 of 53 instances, and 25 of these selections keep the target lower bound on the fresh bank. Over all instances, its paired difference from the original procedure is -0.3 [−5.5, 5.0] points. The capped variant selects a plan on 42 of 53, with 40 retained, and its paired difference from the original procedure is 7.7 [3.5, 12.5] points. Capping removes only the reachability obstruction at individual customers; limits from capacity, sequencing and finite search remain.

Over all instances (Table 7), the capped variant uses 2.5% more distance and 0.43 more vehicles than the original procedure. At $n = 200$ the extra fleet is substantial. The earlier fleet-matched ablation concerns the uncontracted objective and does not cover this cost.

10.4 Clock-aware approximation

The clock-aware rows test a practical approximation that uses departure times. Exact deterministic retiming checks every reported route, so its success at $\sigma = 0$ is a condition for acceptance and

Table 7 Operating cost of the capped variant relative to the original procedure on all 53 instances. Percentage changes are paired by instance; fleet changes are absolute. Capping can lead to different routes and resources, so its reliability gain is not attributed to schedule structure alone.

n	Distance	Scheduled time	Simulated time	Fleet
20	0.0 [0.0, 0.0]	0.0 [0.0, 0.0]	0.0 [0.0, 0.0]	0.00 [0.00, 0.00]
50	2.2 [0.0, 5.4]	4.6 [0.0, 11.4]	4.5 [0.0, 11.3]	0.33 [0.00, 0.87]
100	4.9 [1.1, 8.9]	6.8 [1.7, 12.1]	6.6 [1.7, 11.8]	0.30 [-0.20, 0.80]
200	6.3 [1.6, 10.9]	10.5 [3.1, 18.0]	10.0 [2.9, 17.1]	1.88 [0.50, 3.38]

Table 8 Clock-aware approximation. Each row uses the same instances for every column: those where the approximation found an exact deterministic schedule. Probabilities are percentages, and paired changes are percentage points. The two matrix solves share a five-second search budget.

n	Found	Nominal	Clock-aware	Original P	Clock minus N
20	19/20	21.2 [18.6, 24.0]	39.5 [37.9, 41.3]	84.1 [73.2, 93.6]	18.4 [15.1, 21.4]
50	14/15	17.5 [16.0, 19.3]	36.5 [35.4, 37.7]	74.8 [60.8, 87.1]	19.0 [17.2, 20.9]
100	10/10	25.4 [23.9, 26.8]	44.5 [43.0, 46.0]	81.3 [68.8, 93.2]	19.1 [17.6, 20.8]
200	6/8	26.6 [24.5, 28.9]	43.7 [42.3, 45.3]	81.7 [68.7, 93.6]	17.2 [15.7, 18.7]

says nothing about robustness to random delay. Conservative-speed planning also guarantees deterministic on-time service within the planned horizon, because its planned arc times are never shorter than the deterministic ones. Both 100% results hold by construction.

In Table 8, the approximation exceeds nominal distance by 17.2–19.1 points on the instances where it finds a schedule, but stays below the original procedure. The latest-feasible-departure rule and the two-pass search are part of this result, and the comparison does not show what an optimal time-dependent planner could achieve.

10.5 Matched disturbance sensitivity

Table 9 Fixed-plan sensitivity on the 46 common instances. Each cell uses 1,000 common scenarios from a bank separate from the main validation bank; plans are not re-screened or re-optimized for each setting. C, P and SC are conservative speed, the original procedure and screened conservative speed.

σ	ρ	C (%)	P (%)	SC (%)	P minus C (points)
0.10	0.0	91.8 [88.2, 95.0]	88.8 [81.3, 95.3]	97.1 [94.3, 99.2]	-3.1 [-10.7, 3.9]
0.10	0.5	84.6 [81.5, 87.4]	87.8 [81.5, 93.5]	93.2 [90.1, 95.9]	3.2 [-3.3, 9.3]
0.10	0.9	81.2 [78.3, 84.0]	87.6 [81.4, 93.3]	91.5 [88.2, 94.5]	6.4 [0.0, 12.5]
0.20	0.0	69.8 [64.1, 75.5]	81.0 [71.4, 89.7]	86.7 [81.0, 92.0]	11.2 [1.4, 20.6]
0.20	0.5	65.6 [63.2, 67.9]	82.8 [75.9, 89.2]	82.9 [78.0, 87.5]	17.2 [10.3, 23.7]
0.20	0.9	67.2 [65.3, 68.9]	82.5 [76.8, 87.9]	82.1 [77.9, 86.2]	15.4 [9.5, 20.9]
0.30	0.0	41.7 [37.4, 46.0]	72.1 [61.1, 82.4]	69.4 [60.6, 77.8]	30.5 [19.8, 40.5]
0.30	0.5	54.4 [52.5, 56.3]	74.0 [67.8, 79.8]	72.1 [67.1, 77.0]	19.6 [13.6, 25.3]
0.30	0.9	60.6 [59.3, 61.9]	75.6 [70.7, 80.1]	74.4 [70.6, 78.2]	15.0 [10.0, 19.6]

The comparison of interest in Table 9 is the paired difference P minus C, which Fig. 1 plots across the disturbance grid. It shows how the same fixed plans compare under each setting. This column alone cannot separate screening from the contracted candidate family, and the nine cells share instances and common random numbers.

For independent disturbances, raising σ from 0.10 to 0.30 moves the paired difference from -3.1 [-10.7, 3.9] to 30.5 [19.8, 40.5] points. This supports a growing advantage of the complete

screened procedure in that setting. At $\rho = 0.9$ the increase is not monotonic, so the data do not show that the value of screening always rises with the noise scale.

10.6 Diagnosis of screening failures

Of the 21 instances on which the original procedure admitted no plan, 17 reach a buffer at which some customer is provably unreachable, and 4 have no such certificate up to 60 minutes. On 8 instances the solver also returned no plan at a buffer not covered by a certificate. Those outcomes remain unresolved: failing a necessary reachability test proves infeasibility, but passing it does not show that a feasible multi-customer routing exists. Appendix B.1 gives the audit for each instance.

A further exploratory check took the 8 failures whose search ended with no plan and no reachability certificate, and solved exactly that contracted model again with a longer budget: 30 s, then 120 s for any still without a plan. In 4 cases a plan was found at that contraction level, so the original failure was a search limit. The remaining 4 still had no plan after 120 s, which does not prove infeasibility. The script and records are in the reproducibility archive [5].




Fig. 1 Paired difference in service probability between the original procedure (P) and conservative-speed planning (C) on the 46 common instances. Circles, squares and triangles mark $\rho = 0$, 0.5 and 0.9, and the dashed line marks zero. Faint vertical lines mark $\sigma = 0.10$, 0.20 and 0.30; markers are offset slightly around them only for legibility. Error bars are 95% bootstrap intervals over instances. Plans are fixed across this disturbance sweep, and the intervals are not simultaneous across cells

10.7 First confirmatory study

The four hypotheses of the prespecified protocol were tested on 53 new instances (Table 10). Every decision uses the unrounded interval limits, and the protocol’s validation checks passed before any quantity was computed.

All four hypotheses were confirmed. The original procedure exceeded conservative speed by 15.7 points, and capping added 11.7 points. Capping admitted 46 instances against 32 and lost none that the original procedure admitted (14 gained, 0 lost). For independent disturbances, the advantage over conservative speed was 28.3 points larger at $\sigma = 0.30$ than at $\sigma = 0.10$. Of

Table 10 First confirmatory study on 53 new instances. Differences are in percentage points of validated service-event probability; intervals are paired bootstrap intervals over instances. Conservative speed returned a plan on 50 of 53 instances, and H1 and H4 use those instances, as the protocol specifies.

Hypothesis	Quantity	Estimate [interval]	Level	Decision
H1 (primary)	Original procedure minus conservative speed, 50 instances	15.7 [7.5, 23.6]	97.5%	confirmed
H2 (primary)	Capped minus original procedure, 53 instances	11.7 [6.0, 17.9]	97.5%	confirmed
H3	Admitted by capped / original procedure	46 / 32 (capped only 14, original only 0)	exact $p = 0.00012$	confirmed
H4	Advantage at $\sigma = 0.30$ minus at $\sigma = 0.10$ ($\rho = 0$), 50 instances	28.3 [19.6, 36.4]	95%	confirmed

the plans admitted on the screening bank, 31 of 32 (original) and 45 of 46 (capped) kept the 0.95 lower bound on the fresh bank.

Averaged over all instances, service probability was 79.4% for the original procedure, 82.4% for screened conservative speed and 91.0% for the capped procedure, so on average the 95% target is still not met. Screened conservative speed was not a prespecified comparison, and its value is descriptive only. As in the development study, the shortfall comes from instances without an admitted plan. Plans admitted by screening averaged 97.4% for the original procedure and 97.5% for the capped procedure, while the fallback plans averaged 51.9% (21 instances) and 48.7% (7 instances).

H2 and H3 are not independent. When neither procedure admits a plan, both return the same fallback, and 11.6 of the 11.7 points of H2 come from the 14 instances that only the capped procedure admits. The two hypotheses measure one finding twice.

Post hoc fallback. The protocol fixes the fallback: when no plan is admitted, both procedures return the uncontracted slack-aware plan. Returning the conservative-speed plan instead, where that plan exists, would have raised pooled service on the same fresh bank from 79.4% to 85.0% for the original procedure (18 instances changed) and from 91.0% to 93.0% for the capped procedure (6 instances changed). This rule uses no validation outcome, but it was chosen after the results were seen, so the comparison is exploratory. The second confirmatory study tests it prospectively (Section 10.10).

These instances come from the same generator and the same assumed disturbance model as the development study, so the result shows repeatability within that setting. It says nothing about real road networks, calibrated parameters or other instance families.

10.8 Sensitivity to the objective weights

The original procedure was rerun on the 53 development instances with four weight vectors (Table 11); neither confirmatory study varied the weights. Each rerun regenerates the candidate family, screens it on the same bank and validates the selected plan or fallback on the same fresh bank. Conservative-speed planning does not use the weights, so its archived outcome is reused. The rerun with the original weights differs from the archived run by -1.4 points pooled ($[-5.1, 2.2]$) and by one admitted instance (31 against 32). This is the run-to-run variation of a time-limited search.

Swapping the two largest weights changed nothing beyond the run-to-run variation. The distance-heavy vector gave pooled service 3.2 points lower, with a paired interval that includes

Table 11 Weight sensitivity on the 53 development instances. Weights are ordered distance cost, slack deficit, driver duration, fleet use. The advantage over conservative speed uses the 46 instances where conservative speed returned a plan. The last column is the paired difference in service probability from the rerun with the original weights, over all 53 instances (percentage points).

Vector	Weights	Admitted	Pooled (%)	Advantage over conservative [95%]	Paired difference [95%]
Original (rerun)	(0.5693, 0.2643, 0.1055, 0.0609)	31	77.5	16.2 [9.2, 22.8]	–
Slack first	(0.2643, 0.5693, 0.1055, 0.0609)	30	77.4	16.0 [9.6, 22.1]	–0.1 [–3.9, 3.3]
Distance heavy	(0.8, 0.1, 0.05, 0.05)	29	74.3	12.7 [5.3, 19.7]	–3.2 [–7.1, 0.2]
Equal	(0.25, 0.25, 0.25, 0.25)	23	68.7	6.3 [–1.1, 13.7]	–8.8 [–14.3, –3.5]

zero, and its advantage over conservative speed stayed clearly positive (12.7 [5.3, 19.7]). Equal weights changed the most. They admitted 23 instances against 31, lowered pooled service by 8.8 points (paired interval $-8.8 [-14.3, -3.5]$; -15.6 to -2.4 after adjusting the level for three comparisons), and left an advantage over conservative speed whose interval includes zero (6.3 [–1.1, 13.7]). Of the instances whose admission differed, 9 were admitted only under the original weights and 1 only under equal weights (exact $p = 0.021$, unadjusted; the adjusted threshold for three comparisons is 0.017).

This analysis is exploratory. It uses the development instances and four vectors, chosen without searching for the point at which the advantage disappears, and each vector was run once. The records do not store route structure, so the reason equal weights do worse is not tested here. Equal weights give fleet size and driver duration 0.25 each, against 0.06 and 0.11 in the original vector.

10.9 Solomon instances

The unchanged pipeline was run on eight 100-customer Solomon instances, with the study’s time-dependent profile and disturbance model added (Section 8). Within its five-second budget, OR-Tools returned a nominal plan for five of the eight instances (Table 12). PyVRP/HGS finds a feasible plan for the other three, R101, R102 and RC105, within 40 seconds, so these failures are search limits of the generator, not infeasibility of the mapped instances. The failure policy records them as missing plans, and every configuration that depends on them as not attempted.

Table 12 Solomon instances: validated service-event probability (%). A dash marks a configuration with no plan. *Plan* states whether OR-Tools returned a nominal plan; *Admitted* states whether screening admitted a plan for the original procedure, which otherwise returns the slack-aware fallback.

Instance	Plan	Nominal	Conservative	Slack-aware	Original	Capped	Admitted
R101	no	–	–	–	–	–	no
R102	no	–	–	–	–	–	no
R103	yes	14.5	–	45.1	45.1	45.1	no
R112	yes	29.4	78.0	53.7	99.4	99.4	yes
RC104	yes	22.8	82.0	68.9	97.7	97.7	yes
RC105	no	–	–	–	–	–	no
RC106	yes	19.9	65.0	42.1	42.1	42.1	no
RC108	yes	28.3	80.7	51.5	98.8	98.8	yes

Of the five instances with a plan, screening admitted a plan on three (R112, RC104 and RC108). Their validated service probability was 97.7 to 99.4%, against 78.0 to 82.0% for conservative speed and 22.8 to 29.4% for nominal distance. On the other two, the procedure returned the slack-aware fallback: R103 (45.1%; conservative speed returned no plan there, although HGS finds a feasible plan on its matrix) and RC106 (42.1%), where it fell 22.9 points below conservative speed (65.0%). Capping changed no result on any instance.

On RC106, screening admitted no plan, and the slack-aware fallback was weaker than the conservative plan. The conservative-speed fallback examined post hoc in Section 10.7 would have returned 65.0% there.

The eight instances are all the R1 and RC1 files in the source repository. They give no basis for an interval estimate, and no Solomon comparison was prespecified. The admitted plans reached 97.7 to 99.4% and the fallback plans 42.1 to 45.1%, a split similar to that on the synthetic instances. The generator, however, returned no plan on a substantial share of these harder instances.

10.10 Second confirmatory study

In the exploratory run on the 53 development instances, the HGS capped procedure with conservative fallback exceeded HGS conservative-speed planning by 27.0 points, with pooled service of 94.2%. Table 13 gives the prespecified results on the 53 new instances of the second study.

Table 13 Second confirmatory study on 53 new instances. Differences are in percentage points of validated service-event probability; intervals are paired bootstrap intervals over instances. G1 uses the 49 instances where HGS conservative-speed planning returned a plan.

Hypothesis	Quantity	Estimate [interval]	Level	Decision
G1 (primary)	HGS capped with conservative fallback minus HGS conservative, 49 instances	27.9 [23.0, 32.8]	97.5%	confirmed
G2 (primary)	OR-Tools capped, conservative minus slack-aware fallback, 53 instances	0.4 [−0.7, 2.1]	97.5%	not confirmed
G3	OR-Tools uncapped, conservative minus slack-aware fallback, 53 instances	3.8 [1.6, 6.2]	95%	confirmed
G4	HGS minus OR-Tools, both capped with conservative fallback, 53 instances	−0.2 [−4.0, 3.7]	95%	estimate only

G1 was confirmed. With the HGS generator, screening with capped contraction exceeded conservative padding by 27.9 points, and pooled service probability reached 91.5% against 58.4% for HGS conservative speed. Capping was essential with this generator: at 100 and 200 customers the HGS capped family was admitted on 18 of 18 instances, and the uncapped family on 3.

G2 was not confirmed. The OR-Tools capped procedure admitted a plan on 45 of 53 instances and fell back on 8. On 4 of those no conservative-speed plan existed, so both rules returned the same plan. On the remaining 4 the conservative fallback changed service probability by +5.4, −3.1, −12.0 and +30.0 points. The post hoc gain seen for the capped procedure in the first study did not replicate. For the uncapped procedure, which falls back more often, the conservative fallback helped (G3, 3.8 points).

Under the capped procedure with conservative fallback, the two generators differed by −0.2 points (G4), with a 95% interval from −4.0 to 3.7. No equivalence margin was prespecified, so this interval bounds the difference but does not establish equivalence. The HGS $\beta = 0$ plan,

which lacks the slack term, was a weak fallback: the uncapped HGS procedure pooled 51.0% with it and 74.7% with the conservative fallback (descriptive).

Appendix B reports the results of both confirmatory studies by size (Tables 14 and 15, descriptive only), customer outcomes on the matched subset (Table 16), absolute planned resources (Table 17), service-threshold sensitivity (Table 18), the instance-level failure audit (Table 19), and the equal-budget deterministic benchmark against PyVRP/HGS (Table 20).

11 Reproducibility and implementation checks

The study archive [5] contains the candidate generators, the scenario-bank code, the test suite, the experiment drivers, the raw JSON records and the scripts that produce the results tables and the figure. The tests cover FIFO behaviour, lognormal moments, common random numbers that do not depend on evaluation order, Wilson calculations, visit and capacity feasibility, and the rule that contracted candidates are scored against the original windows. Each run records the Python, NumPy and OR-Tools versions, the operating system, the solver budget, the seeds and the scenario counts.

The clock convention matters for reproduction. Relative time zero is 06:00, and the evaluator adds 360 minutes before looking up the full-day speed profile; leaving out that offset would evaluate a 00:00–10:00 shift instead. The vectorized evaluator is checked against the original scalar implementation, including period boundaries, customer lateness and shared arc disturbances.

Seeds are fixed. Admission uses seed 100000 plus the instance seed, all validation uses 5000 fresh scenarios from seed 400000 plus the instance seed, and the 1000-scenario disturbance sweep uses seed 500000 plus the instance seed. The revision protocol was fixed before the new candidate configurations were run, but earlier aggregate results had motivated these analyses, so they are an exploratory extension on known synthetic instances. They are neither preregistered nor an external replication, and the fresh scenario draws do not remove the reuse of the instance set. Plans in the sweep are not re-screened when σ or ρ changes.

The first confirmatory protocol (Section 8) was publicly timestamped, with a hash manifest of the code that runs and analyses it, before its instances were solved; its results are in Section 10.7. One pipeline smoke test, on an instance outside the confirmatory set, was run while the protocol was being prepared, and the protocol discloses it as pilot data. The confirmatory study used only frozen scripts. The weight-sensitivity analysis was run after the freeze with a script outside the frozen file list. Its first run stopped on a coding error that affected only that script, and the corrected script produced the reported results. The second confirmatory study (Section 9) has its own frozen protocol and manifest. Its OR-Tools half uses the frozen scripts of the first study, and its HGS half is deterministic given the scaling ranges. The script that writes its tables and text stops unless its figures equal the saved output of the frozen analysis.

All results tables and the figure are regenerated from the retained records. Internal planning buffers are always evaluated against the offered windows, candidate selection and final assessment use independent banks, and solver-quality comparisons use equal wall-clock budgets on identical deterministic instances. These rules stop diagnostics and within-run selection from using their own validation outcomes.

Use of generative AI tools. Generative AI assistants were used for code development, experiment implementation, code review, analysis checking, and drafting and editing text. They were OpenAI’s ChatGPT (Astra 6) and Anthropic’s Claude (Opus 5, Opus 5.5 and Sonnet 5), and they are not authors. Every experiment was run from the archived code, every reported number regenerates from the archived raw records, and the headline results were recomputed independently of the analysis scripts. Rerunning the solver can return different routes, because each solve stops at a wall-clock limit; Section 10.8 measures this run-to-run variation. The author reviewed all code and text and takes full responsibility for the content.

12 Limitations

The study has eight main limitations.

First, the generators plan on nominal or iteratively frozen matrices, so the procedure can miss better stochastic plans. The OR-Tools generator is also below the state of the art: under an equal five-second budget, its routes are about 11% longer than those of PyVRP/HGS at 200 customers (Table 20). Absolute distance premiums therefore include solver quality, and some uncertified screening failures may be search limits. The second confirmatory study shows that the reliability result does not depend on OR-Tools, but its HGS generator omits the slack-deficit term, so the two generators do not solve identical objectives. The clock-aware approximation is not an exact time-dependent optimizer. Generator and objective choices also differ between screened families, so the experiments are not a full factorial decomposition of every design decision.

Second, the reliability baselines are all based on padding: uniform conservative speed, and the same screening applied to conservative-speed candidates. No robust or chance-constrained VRPTW optimizer [1, 13, 26] is compared.

Third, the disturbance parameters are assumed and have not been calibrated to fleet GPS residuals. The results hold for those values and are not operational forecasts. The exploratory σ - ρ grid does not replace local calibration.

Fourth, the geometry is synthetic. The Solomon check covers only 8 instances with 100 customers, and the generator found plans for 5 of them. The deterministic solver check does not replace Gehring–Homberger benchmarks or a real road network. The confirmatory instances share the generator and disturbance model of the development study, so they test repeatability within that setting.

Fifth, the weights come from one preference vector. Under equal weights the advantage over conservative speed is not distinguishable from zero (Section 10.8), only four vectors were tried, and neither confirmatory study varied the weights.

Sixth, Wilson screening accounts for binomial sampling error for a fixed candidate. Its nominal coverage is not a simultaneous guarantee across an adaptively searched grid, which is why held-out assessment is needed. It also does not prove feasibility if the disturbance distribution is misspecified.

Seventh, the evidence excludes breakdown recovery, heterogeneous fleets, insertion of new orders, rolling-horizon replanning and departure-time recourse. Every vehicle leaves at its planned time in every scenario, which makes the evaluation stricter than an operator who adjusts departures would face.

Eighth, the fallback remains weak. The slack-aware plan can be worse than conservative padding when screening admits nothing (Section 10.9). In the second confirmatory study, the conservative-speed fallback improved the uncapped procedure but did not measurably improve the capped one, which falls back on few instances.

13 Conclusion

The development study separates three mechanisms. Uniformly conservative travel times protect against predictable congestion by construction, but they pad every trip, including off-peak ones. On matched instances the complete screened procedure exceeded them by 8.3–21.4 percentage points on average, although the interval at 200 customers includes zero. For independent disturbances its advantage grew with noise, from -3.1 to 30.5 points. The slack-aware objective without screening fell below conservative padding, so the advantage comes from the complete procedure, in which the contracted candidate family and statistical selection act together. Screening also improves conservative padding: applied to conservative-speed candidates, it performed about as well as the original procedure, and the capped procedure had the highest

service probability of the three (Table 6; exploratory). The fixed 20% pad is therefore the weaker of the two padding comparators.

The first confirmatory study, on 53 new instances with hypotheses fixed in a timestamped protocol before any of them was solved, confirmed all four prespecified hypotheses with the original weight vector. The screened procedure exceeded conservative speed by 15.7 points (97.5% interval 7.5 to 23.6). Capping added 11.7 points (6.0 to 17.9) and admitted 46 instances against 32, losing none. The advantage over conservative speed grew with independent noise. These instances share the generator and the assumed disturbance model of the development study, so the result shows repeatability within that setting.

The advantage depends on the objective weights. Swapping the two largest weights changed nothing beyond run-to-run variation, and a distance-heavy vector kept a clearly positive advantage. With equal weights, fewer instances were admitted (23 against 31), and the advantage over conservative speed fell to 6.3 points, with an interval from -1.1 to 13.7 . Neither confirmatory study varied the weights.

On 8 Solomon instances with the same overlay, OR-Tools found a plan for 5, and PyVRP/HGS found feasible plans for the other 3. Screening admitted 3 of the 5. The admitted plans reached service probabilities near 98%, and the fallback plans stayed below 46%. This exploratory check covers few instances, and it shows that the candidate generator limits what the pipeline can do on harder benchmarks.

The second confirmatory study replaced OR-Tools with hybrid genetic search. With that generator, screening with capped contraction exceeded conservative-speed planning by 27.9 points (97.5% interval 23.0 to 32.8), and pooled service reached 91.5% against 58.4% for conservative speed. The HGS and OR-Tools procedures differed by -0.2 points, with a 95% interval from -4.0 to 3.7 . No equivalence margin was set, so this interval bounds the difference but does not show equivalence. The reliability gain is therefore not an artefact of a weak generator. Capping mattered more with HGS: at 100 and 200 customers, its uncapped family was admitted on 3 of 18 instances and its capped family on 18.

The fallback is the weakest part of the procedure. When screening admits nothing, the original protocol returns the uncontracted slack-aware plan, which on RC106 was far worse than conservative speed (Section 10.9). In the first confirmatory study, a conservative-speed fallback examined post hoc would have raised pooled service to 85.0% for the original procedure and 93.0% for the capped procedure. Tested prospectively in the second study, this rule did not measurably improve the capped procedure (G2, 0.4 points, interval -0.7 to 2.1). That procedure fell back on only 8 instances, the rule changed the plan on 4 of them, and the effects went both ways. The rule did improve the uncapped procedure, by 3.8 points (95% interval 1.6 to 6.2).

Screening failures can be diagnosed. Of the 21 development instances without an admitted plan, 17 meet a certified reachability obstruction at one or more tested contraction levels. A certificate proves that one particular contracted model is infeasible; it does not show that reachability alone made the procedure miss its target. Capping each customer’s contraction at its direct-reachability limit raised admission from 32 to 42 instances. Of the instances with 100 or 200 customers, 18/18 were admitted and 16/18 kept the target lower bound under fresh validation, with mean service probability 97.5–97.6%. The gain costs 2.5% more distance and 0.43 more vehicles on average, rising to 1.88 more vehicles at 200 customers. The earlier fleet-matched experiment concerns the uncontracted objective and does not isolate the mechanism of the capped variant. At 20 customers capping changes nothing. Only two of the six failures there meet a certificate; of the other four, a longer search budget finds a plan for two, and two remain unresolved.

Averaged over all instances, no procedure reaches the 95% target, so both the procedure averages and the compliance of selected plans must be reported. In the development study, plans admitted by screening averaged 97.3% for the original procedure and 97.4% for the capped procedure, while the fallback plans used on the other instances averaged 50.9% and 45.2%. All

11 capped-procedure fallbacks occur at 20 or 50 customers, so closing the remaining gap depends on admitting plans for small instances. In the first confirmatory study the capped procedure reached 91.0% pooled, against 79.4% for the original procedure.

These results depend on assumed disturbance parameters and synthetic geometry. The development findings were exploratory, and the two confirmatory studies support the main ones. Evidence from real road networks, other benchmark families and calibrated disturbance parameters is still needed.

Acknowledgements

The author thanks the researchers who answered methodological questions by e-mail during the project. These discussions informed experimental design; all claims in the paper rest on the reported model, code, and experiments.

Statements and Declarations

Competing interests. The author has no relevant financial or non-financial interests to disclose.

Funding. No funding was received for conducting this study.

Data and code availability. The reproducibility archive contains source code, dependency versions, tests, seeds, raw JSON records, analysis scripts, generated tables and figures, and both confirmatory protocols. It is archived on Zenodo [5]. The two confirmatory protocol freezes, each archived before any instance of its study was solved, are separate records [3, 4]. The Solomon instance files in the archive are third-party benchmark data [15, 19] and remain under their original terms.

Use of generative AI. Described in Section 11.

Ethics approval. Not applicable: the study uses synthetic data only.

Consent. Not applicable.

Author contributions. S.A. is the sole author and conceived the study, implemented the software, ran the experiments, and wrote the manuscript.

A AHP pairwise-comparison matrix

Criteria are ordered as distance cost, slack deficit, driver duration and fleet use. Entry a_{kl} is the judged importance of criterion k relative to l on Saaty’s 1–9 scale:

$$\mathbf{A} = \begin{pmatrix} 1 & 3 & 5 & 7 \\ 1/3 & 1 & 3 & 5 \\ 1/5 & 1/3 & 1 & 2 \\ 1/7 & 1/5 & 1/2 & 1 \end{pmatrix}.$$

Its principal eigenvalue is $\lambda_{\max} = 4.0685$, and the normalised principal eigenvector is

$$\mathbf{w} = (0.5693, 0.2643, 0.1055, 0.0609).$$

With Saaty’s random index 0.90 for four criteria, the consistency index is $(4.0685 - 4)/3 = 0.0228$ and the consistency ratio is 0.0254. As stated in Section 4, the third row and column were set when the third criterion was total distance and were not re-elicited for driver duration.

B Additional outcomes and failure records

Table 14 First confirmatory study by size (descriptive). For each procedure: the number of instances admitted by screening, with the number that kept the 0.95 lower bound on the fresh bank in parentheses, and the mean validated service-event probability (%) over all instances, including fallbacks.

n	Inst.	Original		Screened conservative		Capped	
20	20	17/20 (17)	90.2	15/20 (15)	87.8	17/20 (17)	90.2
50	15	8/15 (8)	74.6	8/15 (8)	80.3	12/15 (12)	87.4
100	10	4/10 (3)	71.3	3/10 (3)	80.0	10/10 (9)	97.1
200	8	3/8 (3)	71.5	2/8 (2)	75.4	7/8 (7)	92.4

Table 15 Second confirmatory study by size (descriptive): instances admitted by the capped procedure, and mean validated service-event probability (%) over all instances, including fallbacks.

n	Inst.	OR-Tools capped			HGS		
		Admitted	Slack fb.	Cons. fb.	Admitted	Conservative	Capped, cons. fb.
20	20	13	83.8	83.3	14	56.8	85.8
50	15	14	93.7	95.7	13	61.9	92.4
100	10	10	97.9	97.9	10	64.1	97.4
200	8	8	97.5	97.5	8	48.9	97.1

Table 16 Customer outcomes on the instances where conservative-speed planning returned a plan. The clock approximation also excludes the instances where it found no schedule, and the counts show this.

Configuration	k	On time (%)	Expected late
Nominal	46	83.6 [82.5, 84.6]	10.00 [7.95, 12.20]
Conservative speed	46	93.9 [93.4, 94.4]	3.47 [2.84, 4.14]
Slack-aware	46	90.9 [90.0, 91.7]	5.46 [4.38, 6.65]
Fleet-matched	46	83.0 [81.9, 84.0]	10.47 [8.29, 12.87]
Original procedure	46	96.6 [95.1, 97.9]	2.20 [1.26, 3.31]
Screened conservative	46	97.1 [96.3, 97.8]	2.20 [1.47, 2.99]
Capped procedure	46	97.7 [96.4, 98.9]	0.91 [0.56, 1.32]
Clock approximation	44	88.0 [87.2, 88.7]	6.73 [5.39, 8.23]

Table 17 Mean absolute planned resources on the same instances as Table 16: distance in km, scheduled duration in minutes, and vehicles. These are not claims of optimality.

Configuration	k	Distance	Duration	Fleet
Nominal	46	1180.7 [989.9, 1381.0]	2906.5 [2379.9, 3481.3]	7.67 [6.46, 8.98]
Conservative speed	46	1216.2 [1021.5, 1422.0]	3352.5 [2738.5, 4011.0]	8.30 [6.93, 9.76]
Slack-aware	46	1196.3 [1003.0, 1404.0]	2902.0 [2380.3, 3466.2]	7.46 [6.22, 8.78]
Fleet-matched	46	1174.2 [989.2, 1372.8]	2833.6 [2330.6, 3382.6]	7.46 [6.22, 8.76]
Original procedure	46	1247.9 [1049.3, 1453.0]	3171.2 [2620.9, 3777.3]	8.17 [6.89, 9.57]
Screened conservative	46	1252.8 [1056.4, 1461.2]	3494.5 [2874.0, 4161.0]	8.76 [7.35, 10.28]
Capped procedure	46	1283.5 [1076.5, 1506.8]	3317.3 [2709.4, 3964.1]	8.52 [7.13, 10.04]
Clock approximation	44	1137.2 [955.5, 1330.9]	2868.3 [2335.0, 3456.8]	7.64 [6.32, 9.11]

Table 18 Sensitivity to the service threshold κ on the 46 common instances, with the same plans and the same fresh validation bank in every column. At $n = 20/50/100/200$, the number of customers allowed to be late is 2/5/10/20 at $\kappa = .90$, 1/2/5/10 at .95, 0/1/2/4 at .98, and zero at 1.00.

κ	N	C	S	P	SC
0.90	38.8 [36.5, 41.2]	80.6 [78.8, 82.4]	68.8 [65.9, 71.7]	89.2 [84.2, 93.7]	91.2 [88.5, 93.8]
0.95	21.9 [20.3, 23.5]	65.6 [63.1, 68.0]	52.8 [49.9, 55.8]	83.1 [76.3, 89.4]	82.8 [77.9, 87.6]
0.98	11.2 [10.3, 12.2]	48.2 [45.6, 50.8]	36.4 [33.8, 39.0]	75.9 [68.2, 83.3]	73.2 [67.1, 79.3]
1.00	6.4 [5.4, 7.4]	34.4 [31.8, 37.0]	26.1 [23.8, 28.5]	68.2 [59.6, 76.2]	62.4 [53.8, 70.9]

B.1 Instance-level failure audit

Table 19 Every instance on which the original procedure admitted no plan. Best β and the lower confidence bound (LCB) refer to the best returned candidate; the no-plan β is the first buffer at which the solver returned no plan; the proof β is the first buffer with a reachability certificate. Unresolved marks at least one no-plan result without a certificate. A dash means none up to 60 minutes.

n	Seed	Best β	LCB (%)	No-plan β	Proof β	Unresolved
20	7200	45	90.7	50	–	Yes
20	7201	35	93.5	40	–	Yes
20	7210	35	90.8	45	–	Yes
20	7212	5	37.9	10	–	Yes
20	7216	55	90.7	60	60	No
20	7218	10	65.9	15	50	Yes
50	7501	25	87.5	30	35	Yes
50	7503	10	64.1	15	15	No
50	7504	35	91.3	40	40	No
50	7506	0	44.7	5	5	No
50	7508	35	85.9	40	40	No
50	7511	35	88.5	40	50	Yes
50	7513	40	93.6	45	60	Yes
100	8003	20	85.3	25	25	No
100	8004	15	81.5	20	20	No
100	8007	5	73.6	15	15	No
100	8008	25	91.7	35	35	No
200	9004	10	71.2	15	15	No
200	9005	20	88.9	25	25	No
200	9006	25	94.9	40	40	No
200	9007	20	86.4	25	25	No

Table 20 Retained equal-budget benchmark on the deterministic core: excess distance of OR-Tools routes over PyVRP/HGS routes (%). Counts are valid pairs at the one-second and five-second budgets; one pair is missing at one second.

n	Valid pairs (1s/5s)	One second	Five seconds
20	19/20	1.0 [0.3, 1.8]	0.0 [0.0, 0.0]
50	15/15	4.8 [3.3, 6.2]	3.0 [1.9, 4.2]
100	10/10	8.5 [5.9, 11.1]	6.3 [4.7, 7.9]
200	8/8	32.9 [29.4, 36.2]	11.1 [9.3, 13.5]

References

- [1] Agra A, Christiansen M, Figueiredo R, et al (2013) The robust vehicle routing problem with time windows. *Comput Oper Res* 40(3):856–866. <https://doi.org/10.1016/j.cor.2012.10.002>
- [2] Ahmed S, Shapiro A (2008) Solving chance-constrained stochastic programs via sampling and integer programming. In: Chen ZL, Raghavan S (eds) *State-of-the-Art Decision-Making Tools in the Information-Intensive Age*. *TutORials in Operations Research*, INFORMS, Hanover, MD, p 261–269, <https://doi.org/10.1287/educ.1080.0048>
- [3] Ali S (2026) Reliability-screened routing: research code and evidence archive. Zenodo, <https://doi.org/10.5281/zenodo.23185705>, confirmatory protocol freeze, commit 41fae67d9520bd5dd01fcac87e680a5a69325d10
- [4] Ali S (2026) Reliability-screened routing: research code and evidence archive. Zenodo, <https://doi.org/10.5281/zenodo.23208320>, confirmatory protocol 2 freeze, commit 3ba6b71288a2a51872ab6a877fa86eb28ff7671d
- [5] Ali S (2026) Reliability-screened routing: research code and evidence archive. Zenodo, <https://doi.org/10.5281/zenodo.23210189>, version v2.0, software and dataset
- [6] Bomboi F, Buchheim C, Prunte J (2022) On the stochastic vehicle routing problem with time windows, correlated travel times, and time dependency. *4OR* 20(2):217–239. <https://doi.org/10.1007/s10288-021-00476-z>
- [7] Brown LD, Cai TT, DasGupta A (2001) Interval estimation for a binomial proportion. *Stat Sci* 16(2):101–133. <https://doi.org/10.1214/ss/1009213286>
- [8] Ehmke JF, Campbell AM, Urban TL (2015) Ensuring service levels in routing problems with time windows and stochastic travel times. *Eur J Oper Res* 240(2):539–550. <https://doi.org/10.1016/j.ejor.2014.06.045>
- [9] Errico F, Desaulniers G, Gendreau M, et al (2018) The vehicle routing problem with hard time windows and stochastic service times. *EURO J Transp Logist* 7(3):223–251. <https://doi.org/10.1007/s13676-016-0101-4>
- [10] Gendreau M, Jabali O, Rei W (2016) 50th anniversary invited article—future research directions in stochastic vehicle routing. *Transp Sci* 50(4):1163–1173. <https://doi.org/10.1287/trsc.2016.0709>
- [11] Ghiani G, Guerriero E (2014) A note on the Ichoua, Gendreau, and Potvin travel time model. *Transp Sci* 48(3):458–462. <https://doi.org/10.1287/trsc.2013.0491>
- [12] Ichoua S, Gendreau M, Potvin JY (2003) Vehicle dispatching with time-dependent travel times. *Eur J Oper Res* 144(2):379–396. [https://doi.org/10.1016/S0377-2217\(02\)00147-9](https://doi.org/10.1016/S0377-2217(02)00147-9)
- [13] Jaillet P, Qi J, Sim M (2016) Routing optimization under uncertainty. *Oper Res* 64(1):186–200. <https://doi.org/10.1287/opre.2015.1462>
- [14] Li X, Tian P, Leung SCH (2010) Vehicle routing problems with time windows and stochastic travel and service times: Models and algorithm. *Int J Prod Econ* 125(1):137–145. <https://doi.org/10.1016/j.ijpe.2010.01.013>
- [15] ML4VRP (2024) ML4VRP 2024 competition: instances and evaluator. GitHub, URL <https://github.com/ML4VRP/ML4VRP2024>, instances/cvrptw. Accessed 6 October 2026
- [16] Perron L, Furnon V (2026) OR-Tools. URL <https://developers.google.com/optimization>, version 9.15
- [17] Saaty TL (1980) *The Analytic Hierarchy Process*. McGraw-Hill, New York
- [18] Serrano B, Florio AM, Minner S, et al (2026) Contextual stochastic vehicle routing with time windows. *INFORMS J Comput* <https://doi.org/10.1287/ijoc.2025.1189>

- [19] Solomon MM (1987) Algorithms for the vehicle routing and scheduling problems with time window constraints. *Oper Res* 35(2):254–265. <https://doi.org/10.1287/opre.35.2.254>
- [20] Taş D, Dellaert N, van Woensel T, et al (2014) The time-dependent vehicle routing problem with soft time windows and stochastic travel times. *Transp Res Part C Emerg Technol* 48:66–83. <https://doi.org/10.1016/j.trc.2014.08.007>
- [21] Toth P, Vigo D (eds) (2014) *Vehicle Routing: Problems, Methods, and Applications*, 2nd edn. MOS-SIAM Series on Optimization, SIAM, Philadelphia, <https://doi.org/10.1137/1.9781611973594>
- [22] Vareias AD, Repoussis PP, Tarantilis CD (2019) Assessing customer service reliability in route planning with self-imposed time windows and stochastic travel times. *Transp Sci* 53(1):256–281. <https://doi.org/10.1287/trsc.2017.0748>
- [23] Vidal T (2022) Hybrid genetic search for the CVRP: Open-source implementation and SWAP* neighborhood. *Comput Oper Res* 140:105643. <https://doi.org/10.1016/j.cor.2021.105643>
- [24] Wilson EB (1927) Probable inference, the law of succession, and statistical inference. *J Am Stat Assoc* 22(158):209–212. <https://doi.org/10.1080/01621459.1927.10502953>
- [25] Wouda NA, Lan L, Kool W (2024) PyVRP: A high-performance VRP solver package. *INFORMS J Comput* 36(4):943–955. <https://doi.org/10.1287/ijoc.2023.0055>
- [26] Zhang Y, Baldacci R, Sim M, et al (2019) Routing optimization with time windows under uncertainty. *Math Program* 175(1–2):263–305. <https://doi.org/10.1007/s10107-018-1243-y>